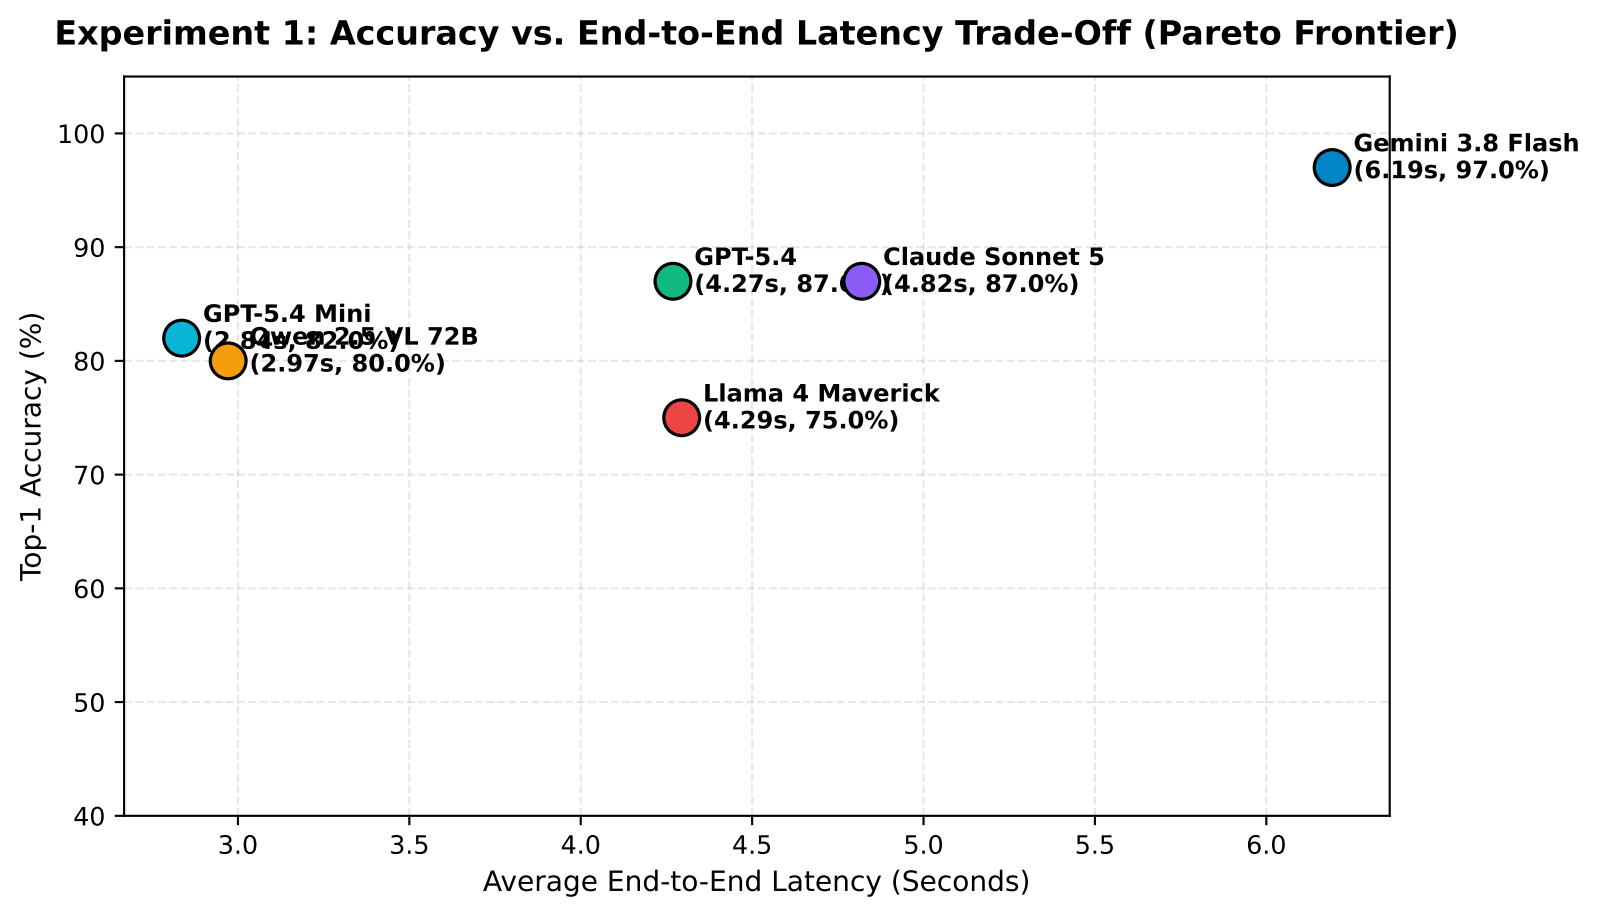
Experiment 1: Accuracy vs. End-to-End Latency Trade-Off (Pareto Frontier)
Top-1 Accuracy (%)
Average End-to-End Latency (Seconds)
100
90
80
70
60
50
40
3.0
3.5
4.0
4.5
5.0
5.5
6.0
Gemini 3.8 Flash
(6.19s, 97.0%)
GPT-5.4
(4.27s, 87.
Claude Sonnet 5
(4.82s, 87.0%)
GPT-5.4 Mini
(2.84s, 82.0%)
Qwen 2.5-VL 72B
(2.97s, 80.0%)
Llama 4 Maverick
(4.29s, 75.0%)