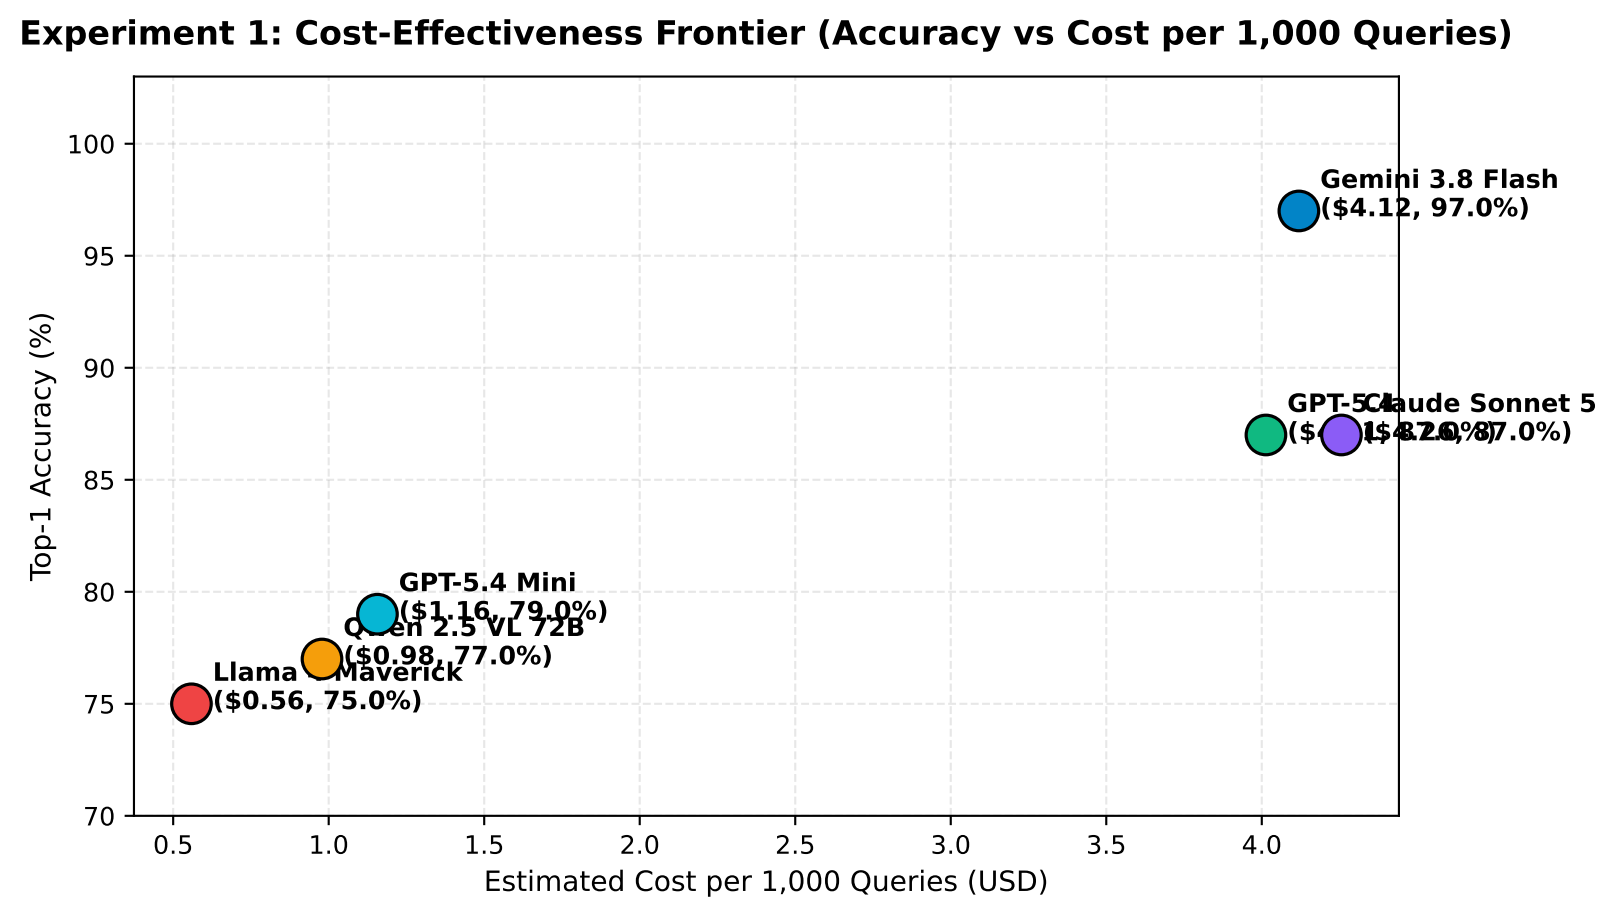
Experiment 1: Cost-Effectiveness Frontier (Accuracy vs Cost per 1,000 Queries)
Top-1 Accuracy (%)
Estimated Cost per 1,000 Queries (USD)
100
95
90
85
80
75
70
0.5
1.0
1.5
2.0
2.5
3.0
3.5
4.0
Gemini 3.8 Flash
($4.12, 97.0%)
GPT-5.4
Claude Sonnet 5
GPT-5.4 Mini
($1.16, 79.0%)
Qwen 2.5 VL 72B
($0.98, 77.0%)
Llama 4 Maverick
($0.56, 75.0%)