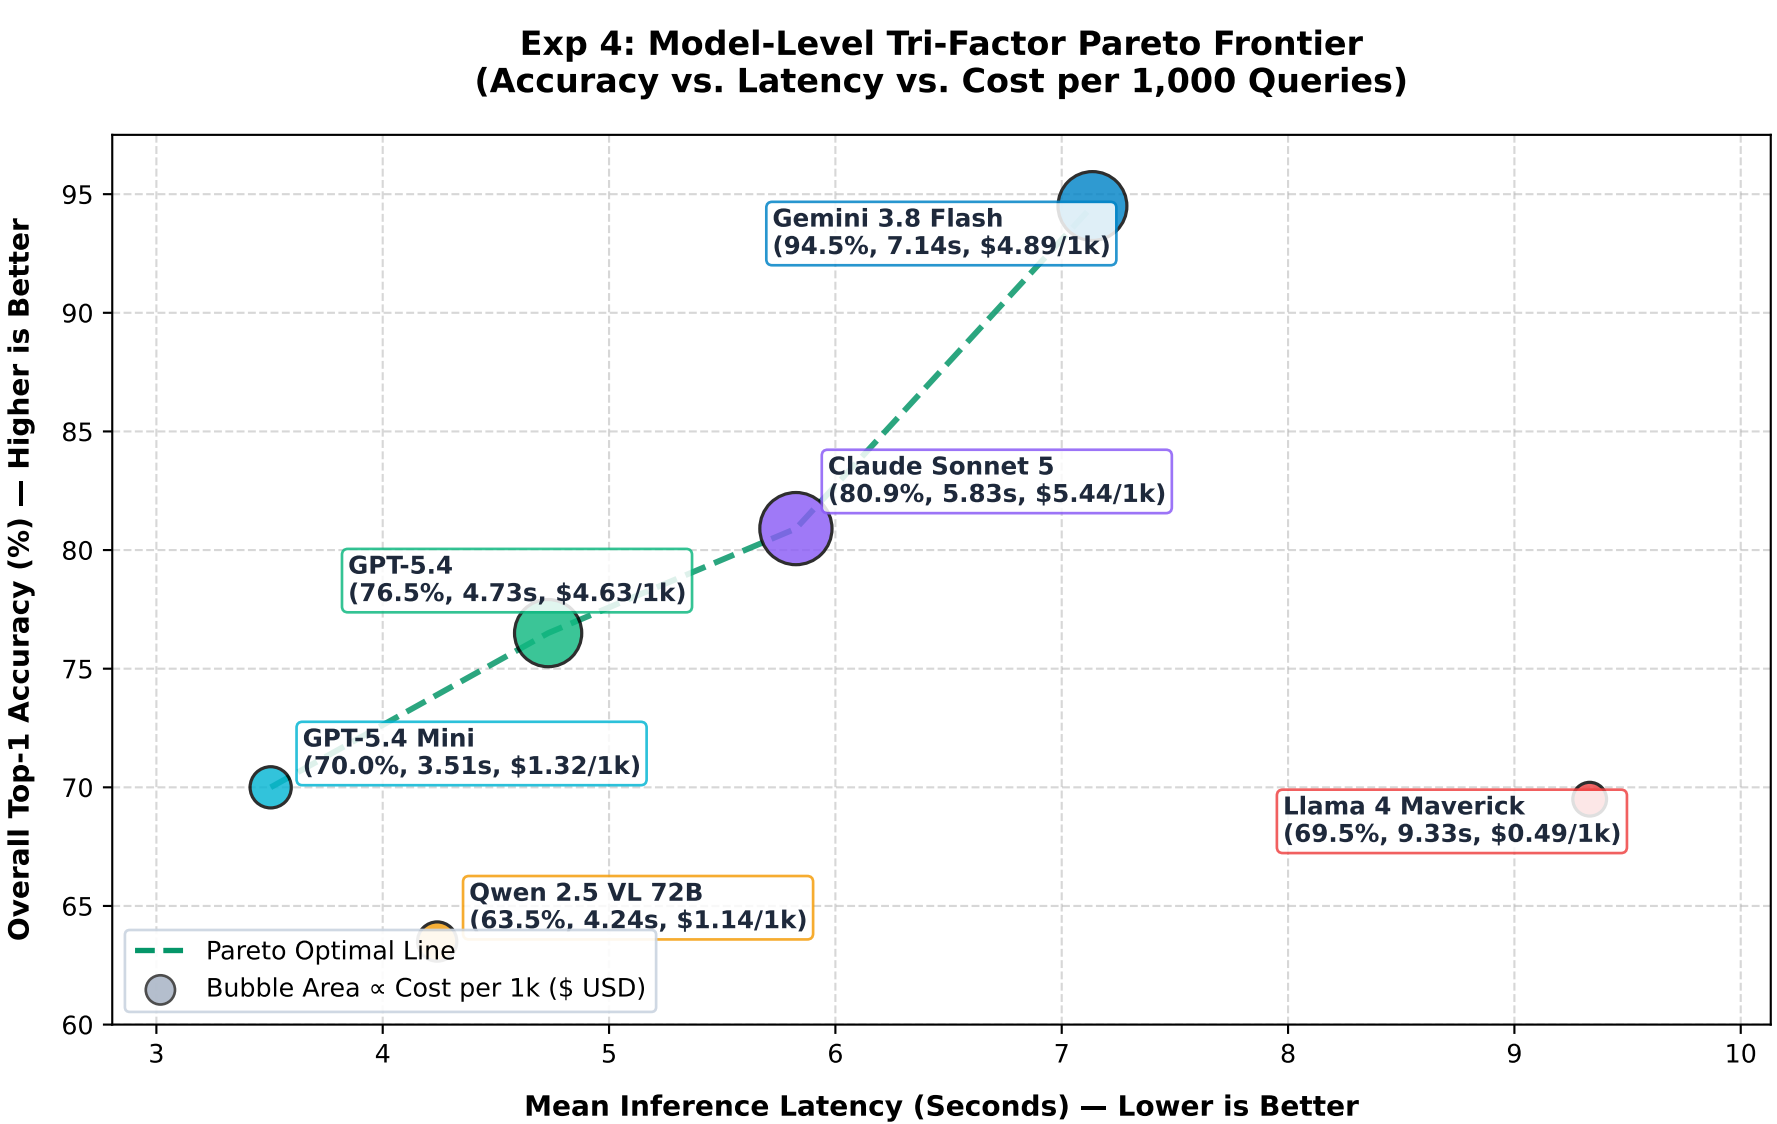
Exp 4: Model-Level Tri-Factor Pareto Frontier
(Accuracy vs. Latency vs. Cost per 1,000 Queries)
Overall Top-1 Accuracy (%) — Higher is Better
Mean Inference Latency (Seconds) — Lower is Better
Gemini 3.8 Flash
(94.5%, 7.14s, $4.89/1k)
Claude Sonnet 5
(80.9%, 5.83s, $5.44/1k)
GPT-5.4
(76.5%, 4.73s, $4.63/1k)
GPT-5.4 Mini
(70.0%, 3.51s, $1.32/1k)
Llama 4 Maverick
(69.5%, 9.33s, $0.49/1k)
Qwen 2.5 VL 72B
(63.5%, 4.24s, $1.14/1k)
Pareto Optimal Line
Bubble Area ∝ Cost per 1k ($ USD)
95
90
85
80
75
70
65
60
3
4
5
6
7
8
9
10